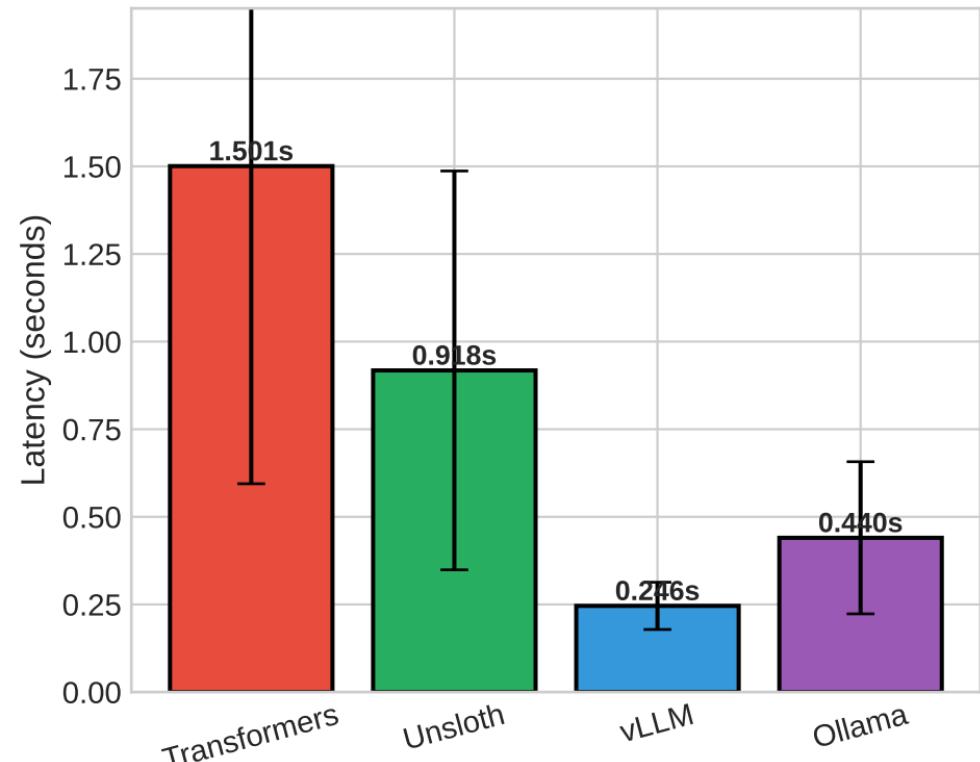
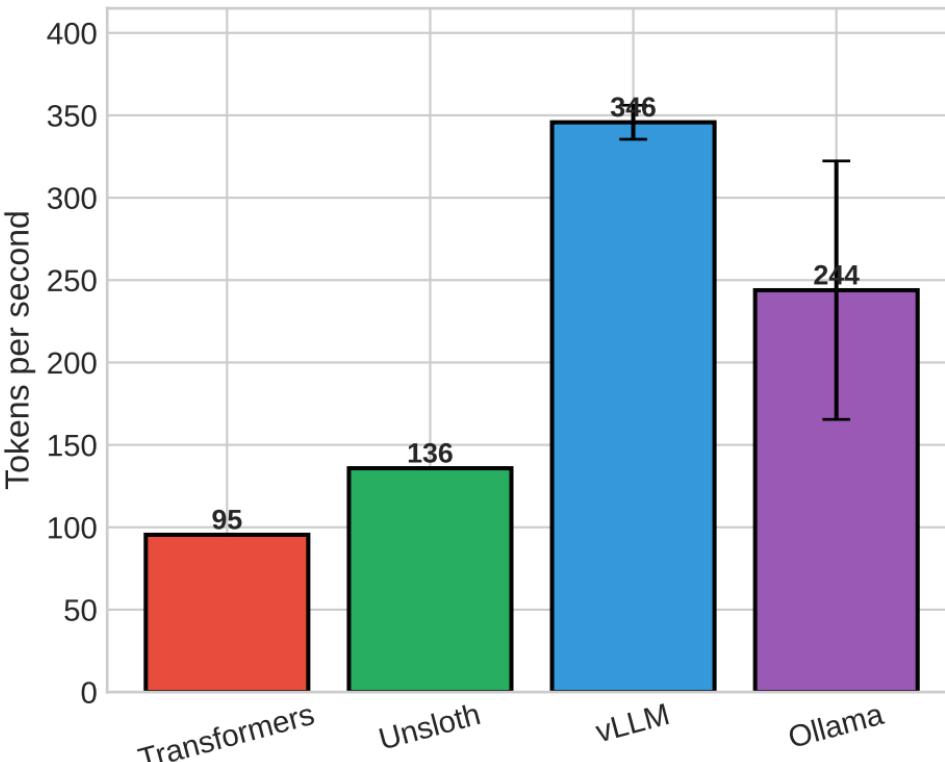


# Inference Benchmark: Framework Comparison

Average Latency  
(lower is better)



Throughput  
(higher is better)



Peak GPU Memory  
(lower is better)

